
Report for Assignment 9: Tabular Q-Learning for Warehouse Navigation

Julian Soh

Department of Mechanical Engineering, Saint Martin's University
ME/EE 467—Machine Intelligence

April 19, 2026

Abstract. Implement Q-learning and State-Action-Reward-State-Action (SARSA) from scratch for the warehouse MDP, compare their learned policies to the value iteration solution from Chapter 4, and investigate the effect of exploration schedules and learning rates on convergence.

1 Introduction

This experiment studies **tabular reinforcement learning** in a small stochastic warehouse navigation problem. The environment is a 4×4 grid world in which an agent started at the lower-left corner and must learn how to reach a goal state at (3, 3) while avoiding a hazardous state at (3, 2). The task was intentionally simple enough that the full state space can be represented exactly with a tabular value function, yet rich enough to highlight important reinforcement-learning ideas such as stochastic transitions, exploration, bootstrapping, and the difference between **off-policy** and **on-policy** learning.

Aside from comparing reward curves, we also compared the learned policies directly against the optimal policy from value iteration. This made it possible to identify which states were learned correctly and where the methods differed, particularly in states in close proximity to the hazard.

Overall, this experiment was designed to show not only whether the agent learned successfully, but also **how different reinforcement-learning algorithms behave under uncertainty**. In particular, it illustrated why Q-learning often learns more aggressive goal-directed behavior, while State-Action-Reward-State-Action (SARSA) can learn safer policies in risky regions because it accounts for the consequences of continued exploration.

2 Methods

2.1 Warehouse and Robot Setup

The warehouse setup and robot behavior are summarized below:

- The warehouse environment is **stochastic**: the intended action is followed with probability 0.8, and each perpendicular action occurs with probability 0.1.
- This transition uncertainty requires the robot to plan under risk, rather than follow a deterministic shortest path.
- Reward design: each non-terminal transition gives a living reward of -0.04 .
- Terminal rewards: reaching the goal gives $+1$, while entering the hazard gives -1 .

- As a result, the learned policy must balance efficiency (shorter routes) against safety (avoiding hazard-adjacent states).

2.2 Comparing Q-learning vs SARSA

The main goal of the experiment was to compare **Q-learning** and **SARSA** on the same control problem (The Warehouse). Both algorithms learn action values from sampled experience, but they differ in a fundamental way.

Q-learning uses the greedy backup target $\max_a Q(s', a)$ and therefore learns toward the optimal policy regardless of the exploratory action actually taken.

SARSA instead updates using the action chosen by the current ϵ -greedy policy, so it learns the value of behaving according to its current exploration strategy. This distinction is especially important near the hazard, where exploratory mistakes can be costly.

We created the accompanying Jupyter notebook to carry out the following sequence of steps:

- Verify required Python packages and imports (NumPy, Matplotlib, and related dependencies).
- Implement the warehouse environment class, including stochastic transitions, rewards, terminal states, and reset/step interaction.
- Implement ϵ -greedy action selection and train Q-learning for 1000 episodes.
- Extract the Q-learning policy using $\pi(s) = \arg \max_a Q(s, a)$ and summarize training rewards.
- Implement and train State-Action-Reward-State-Action (SARSA) for 1000 episodes under the same environment and comparable hyperparameters.
- Extract the SARSA policy and compare it to the Q-learning policy qualitatively on the grid.
- Compute a value-iteration benchmark to obtain the optimal expected return from the start state.
- Plot rolling-average learning curves (window size 50) for Q-learning and SARSA on the same axes, together with the optimal-return reference line.
- Compare learned policies against the optimal policy by counting state-wise matches and visualizing all policies with arrow plots.

2.3 Hyperparameter Sensitivity

The accompanying Jupyter notebook also includes a **hyperparameter sensitivity study** over learning rate α and exploration decay rate ρ to examine how quickly different configurations converge and which settings are unstable or fail to learn reliably.

3 Results and Discussion




Figure 1: Rolling-average learning curves for Q-learning and SARSA, with optimal-return reference.

3.1 Q-Learning

For Q-learning, the learning-curve behavior in Figure 1 shows effective convergence in the tabular setting.

- Early training is noisy due to high exploration (ϵ near 1), producing wide reward variance and occasional hazard visits.
- As ϵ decays, returns improve and the learned behavior becomes increasingly consistent and goal-directed.
- The reward trend approaches the value-iteration benchmark over training, indicating that the off-policy updates are learning a strong control policy.

3.2 SARSA

For State-Action-Reward-State-Action (SARSA), Figure 1 indicates a stable on-policy learning process with competitive final performance.

- Early episodes show high reward variance due to large exploration probability and frequent non-greedy actions.
- As ϵ decays, returns improve and the learned policy stabilizes, indicating convergence in the tabular setting.
- Final performance becomes competitive with Q-learning, while the learning trend is generally smoother during training.

3.3 Policy Comparison

Figure 2 summarizes the learned-policy behavior and trajectory tendencies for Q-learning and State-Action-Reward-State-Action (SARSA) in the stochastic warehouse grid.




Figure 2: Policy comparison between Q-learning and SARSA in the warehouse navigation task.

Q-learning Policy Behavior Q-learning uses the optimistic target $\max_a Q(s', a)$, which tends to favor aggressive action choices in risky regions. In Figure 2, this appears as a tendency to route closer to the hazard when that path is shorter.

- The learned route can prioritize shorter trajectories even when they pass near hazardous states.
- Near (3,2), Q-learning is more likely to choose high-reward but higher-risk actions than SARSA.

SARSA Policy Behavior Because SARSA updates from the action actually taken, it internalizes exploration risk in its value estimates. As shown in Figure 2, this produces a more conservative policy near the hazard.

- SARSA tends to keep a wider safety margin around (3,2) than Q-learning.
- The resulting policy is less aggressive but more risk-aware in high-penalty regions.

3.4 Hyperparameter Sensitivity Experiment

Figure 3 shows the 3×3 learning-curve grid for Q-learning across learning rates $\alpha \in \{0.01, 0.1, 0.5\}$ and exploration-decay rates $\rho \in \{0.99, 0.995, 0.999\}$.

From Figure 3, several trends are clear:

- Faster exploration decay ($\rho = 0.99$) generally produces the quickest reward improvement across tested learning rates.
- Moderate decay ($\rho = 0.995$) remains competitive but typically converges slightly later than $\rho = 0.99$.
- Slow decay ($\rho = 0.999$) tends to underperform within 1000 episodes because prolonged exploration delays policy stabilization.

The learning-rate effect is also visible in Figure 3:

- $\alpha = 0.1$ provides the best balance between convergence speed and stability.




Figure 3: Hyperparameter sensitivity for Q-learning across (α, ρ) combinations.

- $\alpha = 0.5$ can improve quickly for favorable decay settings, but exhibits larger oscillations in some panels.
- $\alpha = 0.01$ is stable but learns more slowly, especially when paired with $\rho = 0.999$.

Overall, the best-performing region in this experiment is around $\alpha \in [0.1, 0.5]$ with $\rho \in [0.99, 0.995]$, while configurations with very slow decay are less effective within the fixed training budget.

4 Conclusion

This experiment demonstrated that tabular reinforcement learning can learn effective navigation behavior in a stochastic warehouse environment with sparse terminal rewards and a hazard penalty. Both Q-learning and State-Action-Reward-State-Action (SARSA) improved substantially over training and produced policies that were broadly aligned with the value-iteration benchmark. The learning-curve results confirmed that both methods are viable in this domain, with Q-learning often showing faster reward gains, while SARSA displayed smoother adaptation under ongoing exploration.

The policy-level comparison highlighted the key behavioral difference between off-policy and on-policy learning. Q-learning, driven by the $\max_a Q(s', a)$ target, tended to prefer shorter and more aggressive routes, including routes that pass closer to the hazard when they provide higher expected return. In contrast, SARSA internalized exploration risk through its on-policy update and therefore learned more conservative trajectories that maintain a wider safety margin near hazardous states. This distinction is important in practical robotics settings where risk tolerance and safety constraints matter as much as path efficiency.

The hyperparameter sensitivity experiment further showed that convergence quality is strongly dependent on both learning rate and exploration schedule. Faster or moderate exploration decay ($\rho \in [0.99, 0.995]$) produced better

performance within the fixed 1000-episode budget than very slow decay ($\rho = 0.999$), which kept exploration high for too long. A moderate learning rate ($\alpha = 0.1$) provided the most reliable balance between speed and stability, while higher learning rate ($\alpha = 0.5$) could converge quickly but with more oscillatory behavior, and lower learning rate ($\alpha = 0.01$) converged more slowly.

Overall, the results support three conclusions: (1) both algorithms can learn strong policies in this warehouse MDP, (2) Q-learning is typically more reward-aggressive whereas SARSA is more risk-aware near hazards, and (3) careful hyperparameter selection is essential for stable and efficient convergence.

References

- Rico A.R. Picone. *Engineering Artificial Intelligence*. Self-published, 2026.
- Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach, fourth edition*. Pearson, 2021. ISBN 1292401133.